

## Design/implementation

1. Implement artificial delay in `memoryfs_server.py` to model a server/network slow to respond

`memoryfs_server.py`

- Added CL arg. `-delayat` using argument parser.
- Added `request_count` variable
  - fixed scoping issue with `global`
- Modified Get and Put
  - increments `request_count`
  - checks if Nth request (`request_count % delayat`) => sleep for 10 seconds

2. Implement at-least-once semantics for `memoryfs_client.py`

`memoryfs_client.py`

- Created a loop to retry `self.block_server.Get`, `Put`, and `RSM` requests
  - used `try except else` syntax
  - Exception `socket.timeout` => print `SERVER_TIMED_OUT`
  - Exception `ConnectionRefusedError` => print `SERVER_DISCONNECTED`
  - if any exception occurred, it would sleep for `RETRY_INTERVAL`
  - if no exception occurred the loop was broken

3. Implement an in-disk key/value store for `memoryfs_server.py`

`memory_server.py`

- Used `argparser` for CL args:
    - `-initdbm`: an integer flag
      - 1 => initializes zero blocks and writes to dbm
      - 0 => reads from dbm and puts into `RawBlocks`
    - `-dbmfile`: a string that names the file used by dbm in disk
  - updated `Put` and `RSM` to update dbm if `dbmfile` specified
4. Implement a client-side cache for file data blocks read (`Get()`) in the `Read()` function in `memoryfs_client.py`
- Created a `block_cache` dictionary in `Diskblocks`
    - maps `block_number` to `bytearray`
  - Created a `inode_cache` dictionary in `Diskblocks`
    - maps `inode_number` to `gencnt`
  - Incremented `gencnt` in `FileName Unlink` and `Write` functions before `StoreInode()`
  - In `FileName.Read()`
    - If the `gencnt` match
      - If the block is present in cache

- Get blocks from the block\_cache
  - print CACHE\_HIT
- If the block is not present in the cache, Get() from the server and store in the cache
- If the genct do not match
  - invalidate all cache entries for this file inode
  - print CACHE\_INVALIDATED

**concern** if there is a hard link, eg f1 => b0, b1 and f2 => b0 and f2 changes, and f2 is modified, shouldn't the cache be

## Testing Methods

- At-least-once behavior when timeouts and server disconnections occur
  - Timeouts:
    - started server with delayat
    - ensured client request recovered to shell after SERVER\_TIMED\_OUT
  - Server disconnections and persisting data across server restarts:
    - started server with dbm and client
    - created files/directories
    - stopped server
    - made client request
    - restarted server from dbm
    - ensured client request recovered to shell after SERVER\_DISCONNECTED

## Assignment questions

**Q1)** In the code given to you, the Acquire() and Release() calls are placed around operations such as cat and append to ensure they run exclusively in one client at a time. What is one example of a race condition that can happen without the lock? Simulate a race condition in the code (comment out the lock Acquire()/Release() in the cat and append functions, and place sleep statement(s) strategically) to verify, and describe how you did it.

After commenting out the Acquire and Release statement for the append, I placed a sleep statement before self.FileObject.Write on line 211 of memory\_shell\_rpc.py. This allowed me to simulate reading the inode with 2 clients concurrently. Since they then wrote to the same offset, the overwrote one another which is a race condition as that isn't how append's intended functionality.

**Q2)** What happens when you don't store the data in disk using dbm on the server, and terminate/restart the server?

Without using dbm to store data in disk, terminating/restarting the server loses all the files because they were stored in memory and not persistent storage.

**Q3)** What are the changes that were made to the Get() and Put() methods in the client, compared to the HW#3 version of the code?

Compared with the HW3 versions of the code, the Get() and Put() methods were modified to make xmlrpc calls to a server instead of modifying a diskblocks object directly with procedure calls.

**Q4)** At-least-once semantics may at some point give up and return (e.g. perhaps the server is down forever). How would you implement this in the code (you don't need to actually implement; just describe in words)

At-least-once semantics may choose to give up to prevent fate sharing when a server is down forever using either a timeout or request limit. Each request (Get, Put, RSM) can be modified so that the loop for retries doesn't continue if a server hasn't responded after a set time period or exceeds a set request limit.